
Neuro-Symbolic Memory Network (NSMN)

*A Hybrid Embedding–Symbolic Reasoning System for Fact Storage,
Retrieval, and Multi-Hop Inference*

1. System Overview

The Neuro-Symbolic Memory Network (NSMN) is a hybrid reasoning system combining neural embeddings, symbolic information extraction, and multi-hop graph reasoning.

User text input is decomposed into structured triples via a rule-based pattern extractor. Each triple receives both a symbolic representation (subject, predicate, object) and a dense vector embedding generated by a SentenceTransformer model.

Facts are stored simultaneously in:

- (1) a vector database (MemoryStore) supporting cosine similarity retrieval and
- (2) a directed symbolic knowledge graph enabling multi-hop reasoning.

Queries follow a retrieve–compose–verify pipeline:

Neural retrieval identifies relevant facts, a reasoning engine composes candidate multi-step chains using both local symbolic joins and graph shortest-path exploration, a deterministic verifier assigns confidence. The final explanation is rendered by converting chain steps into interpretable natural-language statements.

This project implements a Neuro-Symbolic Memory Network (NSMN)
a unified memory system combining:

1. Neural embeddings for semantic similarity search
2. Symbolic triples (subj–pred–obj) extracted from text
3. Graph-based multi-hop reasoning
4. Local reasoning composition
5. Optional LLM-based augmentation + verification
6. Deterministic scoring + ranking of reasoning chains

The system uses Streamlit for insertion and querying.

Core implementation is spread across:

- *store.py* — Vector memory store
 - *core.py* — Main symbolic-neural reasoning engine
 - *app.py* — Streamlit interface
 - *requirements.txt* — dependencies
-

2. Module-Level Architecture

2.1 MemoryStore (Vector Database Layer)

Purpose

A minimal in-memory vector store enabling

- fact insertion
- dynamic embedding matrix maintenance
- brute-force cosine similarity search

Data Model

Each fact becomes a **MemoryItem** with:

id: UUID

symbol: {subj, pred, obj, raw_text, source}

emb: np.ndarray / None

meta: metadata dict

ts: insertion timestamp

Embedding Matrix Management

MemoryStore maintains an **_emb_matrix**, rebuilt lazily only when needed.

Cosine Similarity Search

search_bruteforce(q_emb, k) implements:

1. L2-normalization of query + matrix
2. Dot-product similarity
3. Argsort for top-k selection

Mathematical formulation:

$$sim(q, x_i) = \frac{q}{\| q \|} \cdot \frac{x_i}{\| x_i \|}.$$

This provides deterministic, model-agnostic ranking.

2.2 SymbolicMemoryNetwork (Core Reasoning Engine)

This is the heart of the system.

It integrates:

- **symbol extraction**
- **embedding generation**
- **graph creation**
- **retrieval**
- **chain composition**
- **verification**
- **LLM augmentation (optional)**

2.2.1 Pattern-Based Symbol Extraction

Method: `_extract_symbols(text)`

Uses ~20 linguistic regex patterns (identity, causation, transformation, inclusion, requirements, etc.) to produce structured triples:

```
{
  pred: "causes" | "part_of" | ...,
  subj: normalized left entity,
  obj: normalized right entity,
  subj_raw: raw text,
  obj_raw: raw text,
  src: "pattern"
}
```

Patterns allow approximate IE without external NLP.

If no pattern matches, a fallback "stmt" triple stores the entire sentence.

2.2.2 Embedding Generation

Method: `_embed(text)`

Uses SentenceTransformer (all-MiniLM-L6-v2) to convert:

"pred subj obj"

→ 384-dim vector.

This ensures retrieval is predicate-aware, not just entity-aware.

2.2.3 Symbol Storage

Method: `add_text(text)`

Pipeline:

1. Extract triples via regex
2. Compute embedding for each triple
3. Store into *MemoryStore*
4. Insert nodes+edges into a directed NetworkX graph
 - o Node = normalized entity
 - o Edge attributes: pred, full metadata

This graph becomes the substrate for multi-hop reasoning.

2.2.4 Retrieval

Method: `retrieve(query, k)`

1. Embed the query string
2. Perform brute-force cosine search in MemoryStore
3. Return top-k triples with their similarity scores

This layer ensures neural recall of semantically related facts.

2.2.5 Multi-Step Reasoning

Method: `compose(retrieved, query, max_hops)`

The reasoning engine constructs **candidate chains** from two mechanisms:

A. Local Symbolic Composition

If:

fact1.obj == fact2.subj

Then chain:

A --p1--> B --p2--> C

Score = mean(similarities of both facts).

This implements basic forward chaining.

B. Graph Multi-Hop Reasoning

Using NetworkX *shortest_simple_paths*, up to *max_hops* edges:

Advantages:

- deterministic
- guarantees simplest explanation first
- avoids cycles
- multi-hop generalization

Chain example:

X → Y → Z → Q

Each hop is converted back to a symbolic triple.

Score heuristic:

$$score = 0.5 + 0.4 \times (\text{chain length})$$

This rewards deeper explanations

C. Single-Fact Fallback

If no multi-hop or compositions, system uses retrieved facts independently.

2.2.6 Chain Verification

Method: *verify_chain(chain, query)*

Two layers:

1. Heuristic Confidence Score

Deterministic:

$$conf = clamp(0.5 + 0.1 \cdot \text{len}(\text{chain}), 0.2, 0.95)$$

Larger chains → higher confidence.

2. Optional LLM Verification

If enabled:

- System passes facts + question to a T5 model
- Asks for 1-sentence verification rationale
- Not used for ranking, only explanation

This avoids LLM hallucination affecting reasoning correctness.

2.2.7 LLM-Based Fact Augmentation (Optional)

Method: `_augment_with_llm(query, retrieved)`

If weak reasoning occurs, the system asks an LLM for:

"bridging facts" that logically connect retrieved facts to the question.

LLM returns JSON describing subject–predicate–object triples, which are inserted back into memory.
This forms a **neuro-symbolic self-extending memory system**.

2.2.8 Final Ranking + Explanation

Method: `answer(query)`

Candidate chains are ranked by:


$$score = 0.4 \cdot raw_score + 0.35 \cdot lexical_overlap(query, chain) + 0.25 \cdot length_bonus$$


Where:


- lexical overlap = token intersection
- length_bonus favors deeper reasoning (academic-style logic)

Then:

- Best chain is rendered to readable English
- UI displays candidate chains + verification info







3. Graph Semantics

The directed graph stores semantic relationships:

node = concept

edge.pred = relation type

edge.meta = provenance

This enables:

- causal chains
- entity linking
- semantic propagation
- conceptual clustering

The system avoids uncontrolled graph explosion by only adding triples from:

- pattern extraction
- LLM-generated facts (optional)
- user input

4. Streamlit Application Layer

The UI provides:

Left Panel

- Text input

- Buttons:
 - **Add to memory** → calls core.add_text()
 - **Ask** → calls core.answer()
- Displays:
 - best chain explanation
 - candidate chains
 - retrieved facts

Right Panel

- Memory browser (list all facts)
 - Clear memory button
-

5. System Strengths

1. Full Neuro-Symbolic Pipeline

This architecture blends:

- neural similarity
- symbolic logic
- graph reasoning
- LLM augmentation

2. Deterministic, Inspectable Reasoning

Unlike LLMs:

- every chain is traceable
- every inference is explainable
- no hallucination distortion

3. Modular & Extensible

Swapping components is trivial:

- embedding model
 - pattern extractor
 - graph constraints
 - reasoning scoring
-

6. Limitations

- pattern dependency,
- no learning mechanism,
- limited linguistic coverage,
- no contradiction detection,
- no probabilistic interpretation,
- performance not optimized for large-scale memory.

7. Future Directions

Neural Extensions

- relation extraction via transformers,
- entity linking across sentences,
- transformer-based paraphrase compression.

Symbolic Enhancements

- graph visualization,
- rule induction,
- contradiction detection.

Hybrid Models

- LLM verification layers,
 - uncertainty scoring,
 - reinforcement for chain ranking.
-

8. Conclusion

The NSMN architecture integrates neural semantic retrieval with explicit symbolic reasoning in a fully inspectable pipeline. By isolating vector search, symbol extraction, graph construction, multi-hop reasoning, and verification into modular components, the system supports deterministic, interpretable inference while remaining extensible to stronger encoders or LLM-based augmentation.

Its end-to-end structure pattern-based extraction, embedding-driven recall, graph-based composition, heuristic verification, and explanation rendering creates a functional neuro-symbolic reasoning loop suitable for research-level applications in explainable AI, factual QA, and hybrid memory architectures.